TempoBench-Grounded Human/AI Benchmark Game Family (GF-01):
Standalone Formal Specification with Decision and Hypothesis Glossary

Bobby Veihman
Project Specification

Last edited: February 15, 2026

Abstract

This document gives a standalone formal specification for the current TempoBench-grounded benchmark family `GF-01`. Ground truth is defined over a canonical substrate: *(HOA/reactive-system object, finite trace, explicit intervention semantics)*. Primary scoring is machine-checkable through structured certificates rather than free-form explanation. To keep the document readable without access to repository-internal notes, all referenced decision labels (DEC-*), hypothesis labels (HYP-*), and renderer labels (GF-01-R*) are explicitly defined here.

Public implementation repository. A publishable mirror of the benchmark implementation artifacts is available at https://github.com/official-Auralin/Benchmark-game.

1 Reading Guide and Label Glossary

1.1 What the labels mean

- **DEC-***: locked project decisions that are normative in this specification.

- **HYP-***: explicit hypotheses (not yet treated as facts), each with a validation plan.

- **Q-***: open or deferred questions.

- **EVAL-***: evaluation-condition tracks (closed-book, tool-augmented, oracle-ceiling).

- **MET-***: metric/scoring regimes.

- **Code/field identifiers**: shown in monospace (e.g., `family_id`, `eval_track`).

1.2 Renderer labels used in this document (see Definition 8)

- **GF-01-R1**: side-scroller renderer where each segment/column corresponds to one timestep; interventions for timestep t are chosen before advancing to $t + 1$.

- **GF-01-R2**: puzzle-style renderer that presents the same canonical observation information in puzzle-board form.

- **GF-01-R3**: turn-locked grid/tower renderer under consideration; treated as a renderer variant only if semantic parity checks pass.

1.3 Normative decisions referenced in this spec

DEC-001
>   report `EVAL-CB`, `EVAL-TA`, and `EVAL-OC` separately; do not mix tracks.

DEC-001a
>   keep naming separation: `EVAL-*` for conditions, `MET-*` for metric regimes.

DEC-002
>   interventions occur online during play.

DEC-003
>   intervention semantics are ternary per (t, AP): set 0, set 1, unchanged.

DEC-004
>   primary efficiency cost is number of distinct intervened timesteps; atom-count is secondary.

DEC-005
>   action availability must not leak hidden state.

DEC-012
>   normative minimality is singleton-removal.

DEC-012a
>   subset-minimality is optional and non-scoring.

DEC-012b
>   subset-minimality diagnostics must be exact-only when run.

DEC-014
>   hard mode uses exact-time objective; normal mode uses windowed objective.

DEC-014a
>   normal-mode window scales with complexity.

DEC-014b
>   current normal-window coefficients are fixed defaults, pending post-pilot calibration.

DEC-016
>   renderer variants are allowed only under semantic parity.

DEC-018
>   per-timestep intervention cardinality is unrestricted, constrained by global budgets.

DEC-019
>   exact normative checks are always mandatory.

DEC-020
>   optional-diagnostic infeasibility uses hybrid caps $(P_{\text{ref}}, A_{\text{ref}})$.

DEC-021

    initial cap freeze timing is tied to the scheduled pilot checkpoint.

DEC-022

    cap derivation uses a dual-threshold rule.

DEC-023

    current provisional dual-threshold defaults are $p = 0.90$, $k = 3$.

DEC-024

    generator determinism is mandatory for fixed input tuple.

DEC-025

    AP/TS PRF under non-unique minima uses deterministic best-match mapping.

DEC-031

    run artifact schemas are versioned and metadata-rich.

DEC-038

    implementation priority is `GF-01-R1` first.

DEC-038a

    `GF-01-R3` promotion requires pre-registered parity/leakage/reporting gates.

2 Objective and Evidence Anchors

2.1 Objective

The benchmark is designed to evaluate temporal causality reasoning, aligned with TempoBench TCE-style tasks where the model must recover causally relevant interventions over time [3].

2.2 Evidence policy

Nontrivial factual claims in this document are backed by cited sources. Project-policy locks are identified by DEC labels. Statements not yet established empirically are marked as HYP labels.

3 Canonical Substrate and Formal Notation

3.1 Finite-trace benchmark scope

TempoBench defines task objects over automata and finite traces for temporal reasoning and temporal causality [3]. Coenen et al. formalize temporal causality in reactive systems in an ω-trace setting with explicit counterfactual structure [2]. This benchmark intentionally uses a finite-trace operational scope while keeping formal causality checks machine-verifiable.

Definition 1 (System and proposition partition)**.** *Each instance uses a system object A with atomic propositions*

$$AP = AP_{\mathrm{in}} \uplus AP_{\mathrm{out}},$$

where AP_{in} are intervenable input propositions and AP_{out} are observable output propositions.

Definition 2 (Base trace). *Let*

$$\tau = (\tau_0, \ldots, \tau_{T-1}), \qquad T \geq 1,$$

where each $\tau_t : AP_{\text{in}} \to \{0,1\}$ is an input valuation.

Intuition. τ is the default input schedule. The player/agent edits this schedule via interventions (see Definition 3, Definition 4, and Definition 5).

3.2 Intervention semantics

Definition 3 (Intervention atom). *An intervention atom is (t, a, v) with $t \in \{0, \ldots, T-1\}$, $a \in AP_{\text{in}}$, $v \in \{0,1\}$.*

Definition 4 (Certificate). *A certificate is a finite set*

$$C \subseteq \{0, \ldots, T-1\} \times AP_{\text{in}} \times \{0,1\}.$$

Definition 5 (Apply operator). *Define the intervened trace $\tau[C] = \text{Apply}(\tau, C)$ (`Apply`) by*

$$\tau[C]_t(a) = \begin{cases} v, & \text{if } (t, a, v) \in C, \\ \tau_t(a), & \text{otherwise.} \end{cases}$$

Malformed certificates with conflicting assignments to the same (t, a) are rejected by structural validation.

Intuition. "Unchanged" is represented by omission: if no atom targets (t, a), base trace value $\tau_t(a)$ remains.

3.3 Run and canonical observation object

Definition 6 (Run).
$$\rho = \text{Run}(A, \tau[C])$$

denotes the benchmark rollout under the intervened trace (`Run`). For a fixed instance bundle and evaluator version, rollout semantics are replay-deterministic (DEC-024, DEC-031).

Definition 7 (Canonical observation object). *At timestep t, the canonical agent-facing object is*

$$O(s_t) = \{t, y_t, \texttt{effect\_status}_t, b_T^{\text{rem}}, b_A^{\text{rem}}, C_{\leq t}, \text{mode}, t^*\},$$

with hidden internal state omitted. All renderers must be semantics-preserving maps from this canonical object.

4 Instance Definition and Game Interpretation

4.1 What GF-01 means (family-level definition)

Definition 8 (GF-01 benchmark family). *GF-01 is a forward-time, intervention-driven benchmark family in which:*

1. the environment evolves over discrete timesteps $t = 0, \ldots, T-1$,

2. at each timestep, an agent (human or AI) can apply do-style edits to input propositions in AP_{in},

3. the episode objective is to achieve a formal effect predicate e with a certificate that passes machine-checkable causal-validity tests (sufficiency + minimality),

4. scoring is computed from structured artifacts (intervention logs and final certificate), not from natural-language explanation.

Plain-language interpretation. `GF-01` is the core game family name for this benchmark. All renderer variants (e.g., side-scroller, puzzle, turn-locked grid/tower) are just different user interfaces over the same formal game: same state semantics, same intervention rules, same ground truth, and same evaluator.

Definition 9 (GF-01 instance). *A `GF-01` instance is*

$$I = (A, AP_{\text{in}}, AP_{\text{out}}, \tau, T, e, t^*, \text{mode}, w(I), B_T, B_A, \text{meta}).$$

Field meanings.

- e: machine-checkable effect predicate,

- t^*: target timestep for effect objective,

- $\text{mode} \in \{\text{normal}, \text{hard}\}$,

- $w(I)$: trailing window size used in normal mode,

- B_T: primary timestep intervention budget,

- B_A: optional atom-budget field (diagnostic/constraint use).

This object is the formal instance contract used throughout the evaluator (see Definition 9).

4.2 Mode-specific effect objective

$$E_{\text{hard}}(\rho, e, t^*) := [\rho_{t^*} \models e],$$

$$E_{\text{norm}}(\rho, e, t^*, w) := \exists t \in [\max(0, t^* - w), t^*]\, [\rho_t \models e].$$

$$E_I(\rho) = \begin{cases} E_{\text{norm}}(\rho, e, t^*, w(I)), & \text{mode} = \text{normal}, \\ E_{\text{hard}}(\rho, e, t^*), & \text{mode} = \text{hard}. \end{cases}$$

This normal/hard split is locked by `DEC-014`.

4.3 Normal-mode window scaling policy

Current locked default:

$$w(I) = \text{clamp}(w_{\min}, w_{\max}, \text{round}(\alpha_0 + \alpha_T T + \alpha_C z(I))),$$

with provisional defaults (`DEC-014b`): $\alpha_0 = 1$, $\alpha_T = 0$, $\alpha_C = 2$, $w_{\min} = 1$, $w_{\max} = 6$, and equal-weight normalized complexity composite $z(I)$ based on TempoBench-linked structural factors [3].

5 Ground-Truth Causality Contract

Definition 10 (Sufficiency)**.**

$$\mathrm{Suff}_I(C) := \mathbf{1}[E_I(\mathrm{Run}(A, \tau[C])) = 1].$$

Definition 11 (Normative minimality (singleton-removal))**.**

$$\mathrm{Min1}_I(C) := \mathbf{1}\Big[\forall c \in C : E_I(\mathrm{Run}(A, \tau[C \setminus \{c\}])) = 0\Big].$$

Definition 12 (Normative validity)**.**

$$\mathrm{Valid}_I(C) := \mathrm{Suff}_I(C) \wedge \mathrm{Min1}_I(C).$$

Definition 13 (Ground-truth certificate set)**.**

$$\mathcal{M}(I) := \{C \mid \mathrm{Valid}_I(C) = 1\}.$$

Multiple distinct valid minimal certificates are allowed.

Normative policy. `DEC-012` locks singleton-removal minimality as normative scorer behavior.

5.1 Optional stronger diagnostic

Optional (non-scoring) exact subset-minimality diagnostic:

Definition 14 (Optional exact subset-minimality diagnostic)**.**

$$\mathrm{SubsetMin}_I(C) := \mathbf{1}[\nexists C' \subset C \ \mathit{with} \ \mathrm{Suff}_I(C') = 1].$$

By policy (`DEC-012a`, `DEC-012b`), this diagnostic is optional and exact-only when run.

6 Coenen Alignment and Scope Divergence

Coenen et al. define causality with explicit PC1/PC2/PC3 structure, contingencies, and strict-subset/property-level minimality [2]. This benchmark intentionally diverges in normative scoring by using finite intervention certificates with singleton-removal minimality. The divergence is explicit and policy-locked (centered on `DEC-012`), not implicit.

7 Evaluation Tracks and Metric Regimes

7.1 Evaluation-condition tracks

- `EVAL-CB`: closed-book (no external tools beyond model-internal reasoning/scratchpad).

- `EVAL-TA`: tool-augmented with predeclared allowlist and full tool logging.

- `EVAL-OC`: solver-oracle ceiling track, reported separately.

Track separation is mandatory (`DEC-001`, `DEC-001a`).

7.2 Metric regimes

The symbols in this section reuse formal objects from Definition 10, Definition 11, Definition 12, and Definition 13.

MET-C (Certified Causality).

$$\text{Score}_C(I) = \mathbf{1}[\text{Valid}_I(C_a) = 1].$$

where the implementation artifact name is `Score_C`.

MET-M (Multi-objective). For agent certificate C_a:

$$M(I) = \mathbf{1}[\text{Valid}_I(C_a) = 1], \quad G(I) = \mathbf{1}[\text{Suff}_I(C_a) = 1],$$

$$\text{Eff}_t(I) = |\text{TS}(C_a)|, \quad \text{TS}(C) = \{t \mid \exists a, v : (t, a, v) \in C\},$$

$$\text{Eff}_a(I) = |C_a|.$$

Per-instance lexicographic key:

$$\kappa(I) = \big(M(I), G(I), -\text{Eff}_t(I), -\text{Eff}_a(I)\big).$$

with implementation key name `kappa`.

Intuition. The scoring priority is: valid minimal causality first, then effect success, then fewer intervention timesteps.

7.3 AP/TS precision-recall-F1 under non-unique minima

TempoBench reports AP/TS causality metrics [3]. With non-unique $\mathcal{M}(I)$ (Definition 13), deterministic best-match target is (DEC-025):

$$C^*(I, C_a) = \arg \max_{C \in \mathcal{M}(I)} F1_{AP}(C_a, C),$$

then tie-break by $F1_{TS}$, then lower Eff_t, then canonical atom hash order.

7.4 Required reporting dimensions

Official reports must stratify by

$$(\texttt{family\_id}, \texttt{eval\_track}, \texttt{renderer\_track}, \texttt{difficulty\_slice}, \texttt{split\_id}).$$

No cross-track pooling is allowed. This follows evidence that aggregate single scores can hide major failures [4] and that contamination/governance choices can distort benchmark conclusions [1, 7, 8].

8 Generator Contract and Difficulty Scaling

8.1 Generator input object

$$G_{\text{in}} = (A, \tau, AP_{\text{in}}, AP_{\text{out}}, F, \text{seed}, \text{cfg}).$$

where implementation-facing field names include `seed` and `cfg`. Generation is defined for any finite trace length $T \geq 1$. Candidates are accepted only if exact normative checking confirms $\mathcal{M}(I) \neq \emptyset$ (see Definition 13).

8.2 Determinism lock

For fixed $(A, \tau, AP_{\text{in}}, AP_{\text{out}}, F, \text{seed}, \text{cfg}, \text{generator\_version})$, output must be bit-identical (`DEC-024`). No wall-clock or system-entropy randomness is permitted.

8.3 Connection to TempoBench generation evidence

TempoBench describes a formal generation/evaluation pipeline using specification-to-automaton synthesis and causal extraction over generated traces [3]. Repo-side benchmark-runner artifacts provide AP/TS-scored execution/logging behavior [6]. Solver-validated generation as a general methodology is also supported by SATBench-style design [9].

9 Optional-Diagnostic Infeasibility Policy

Normative exact checks are always required. Only optional diagnostics may be marked infeasible (`DEC-019`).

Hybrid policy (`DEC-020`):

- platform/runtime caps P_{ref},

- algorithmic-budget caps A_{ref}.

Initial cap freeze timing is aligned with `DEC-021` checkpoint; derivation uses dual-threshold rule from `DEC-022` with provisional defaults $p = 0.90$, $k = 3$ (`DEC-023`).

10 Partial Observability and No-Query Rule

Hidden state is allowed but not directly exposed. Agents gather information only from passive observations and intervention consequences. Explicit query actions are excluded. The agent-facing channel is the canonical observation object in Definition 7. This aligns with interactive-agent evaluation practices that require explicit environment protocols rather than static one-shot prompts [5, 10].

11 Renderer Policy and Current Priority

11.1 Renderer parity requirement

Renderer variants are allowed only if they preserve canonical $O(s)$ semantics, intervention semantics, and scoring outcomes (see Definition 7 and Definition 12).

11.2 Current lock

By `DEC-038`, implementation priority is:

- first: `GF-01-R1` (side-scroller),

- deferred candidate: `GF-01-R3` (turn-locked grid/tower), promoted only after `DEC-038a` gates pass: semantic parity, leakage-audit parity, and track-separated reporting.

12   Active Hypotheses and Validation Plans

- **HYP-005**: exact subset-minimal diagnostics may be costly at scale. Validation: runtime/cost curves versus instance complexity.

- **HYP-006**: forward-time defensive gameplay can preserve TCE fidelity while improving engagement. Validation: matched human/agent pilot.

- **HYP-007/HYP-008**: complexity-scaled normal windows and current trigger constants are useful but provisional. Validation: post-pilot recalibration audit.

- **HYP-011/HYP-012**: current family ranking and decision matrix are inference-first and require empirical backfill.

- **HYP-013/HYP-014**: hybrid infeasibility caps improve cross-lab comparability when frozen from pilot evidence.

- **HYP-015/HYP-016**: Python-first implementation is sufficient for near-term scale but requires larger stratified sweep confirmation.

- **HYP-017**: turn-locked grid/tower renderer may scale better than vertical side-scroller on high-complexity slices. Validation: matched-seed parity/usability pilot versus R1.

13   Deferred Items and Current Risk Note

- **Q-033 (deferred):** confirm profiling gates on larger stratified sweeps with stronger hardware.

- Optional-diagnostic coverage comparability depends on successful cap freeze execution and transparent run metadata.

14   Summary

This specification locks the current normative benchmark contract while keeping hypotheses and deferred items explicitly labeled. The document is self-contained: all key labels and renderer identifiers used in formal sections are defined in Section 1.

References

[1] Zerui Cheng, Stella Wohnig, Ruchika Gupta, Samiul Alam, Tassallah Abdullahi, Jose Alves Ribeiro, et al. Benchmarking is broken - don't let ai be its own judge. In *NeurIPS Datasets and Benchmarks Track*, 2025. URL https://arxiv.org/abs/2510.07575. source_id: SRC-020.

[2] Norine Coenen, Bernd Finkbeiner, Hadar Frenkel, Christopher Hahn, Niklas Metzger, and Julian Siber. Temporal causality in reactive systems. In *Automated Technology for Verification and Analysis (ATVA)*, volume 13505 of *Lecture Notes in Computer Science*, pages 208–224, 2022. doi: 10.1007/978-3-031-19992-9_13. URL https://doi.org/10.1007/978-3-031-19992-9_13. source_id: SRC-003.

[3] Nikolaus Holzer, William Fishell, Baishakhi Ray, and Mark Santolucito. Mechanics of learned reasoning 1: Tempobench, a benchmark for interpretable deconstruction of reasoning system performance. *arXiv preprint arXiv:2510.27544*, 2025. doi: 10.48550/arXiv.2510.27544. URL https://arxiv.org/abs/2510.27544. source_id: SRC-001.

[4] Percy Liang, Rishi Bommasani, Tony Lee, et al. Holistic evaluation of language models. *arXiv preprint arXiv:2211.09110*, 2022. URL https://arxiv.org/abs/2211.09110. source_id: SRC-017.

[5] Xiao Liu, Hao Yu, Hanchen Zhang, Yifan Xu, Xuanyu Lei, et al. Agentbench: Evaluating llms as agents. In *International Conference on Learning Representations (Datasets and Benchmarks Track)*, 2024. URL https://openreview.net/forum?id=zAdUB0aCTQ. source_id: SRC-008.

[6] nik-hz and contributors. Tempobench github repository. https://github.com/nik-hz/tempobench, 2025. URL https://github.com/nik-hz/tempobench. source_id: SRC-002; audited_commit: 1e4d5c1c4b2931dfa1dbfcb7c19444a4ea318e91.

[7] Weijia Shi, Sewon Min, Maria Lomeli, Chunting Zhou, Margaret Li, et al. Open source data contamination in gpt-4 and other llms. *arXiv preprint arXiv:2310.17589*, 2023. URL https://arxiv.org/abs/2310.17589. source_id: SRC-019.

[8] Soham, Anik Santikary, and Ruoxi Jia. A survey on the impact of data contamination in the evaluation of large language models. *arXiv preprint arXiv:2406.04244*, 2024. URL https://arxiv.org/abs/2406.04244. source_id: SRC-018.

[9] Anjiang Wei, Yuheng Wu, Yingjia Wan, Tarun Suresh, Huanmi Tan, Zhanke Zhou, Sanmi Koyejo, Ke Wang, and Alex Aiken. Satbench: Benchmarking llms' logical reasoning via automated puzzle generation from sat formulas. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, 2025. doi: 10.18653/v1/2025.emnlp-main.1716. URL https://doi.org/10.18653/v1/2025.emnlp-main.1716. source_id: SRC-004.

[10] Shuyan Zhou, Frank F. Xu, Hao Zhu, Xuhui Zhou, et al. Webarena: A realistic web environment for building autonomous agents. *arXiv preprint arXiv:2307.13854*, 2023. URL https://arxiv.org/abs/2307.13854. source_id: SRC-015.